

Team Project

José Ignacio Díez Ruiz 100487766

Carlos Roldán Piñero 100484904

Pablo Vidal Fernández 100483812

2023-01-17

Descriptive Analysis and Preprocessing

Reading the data and setting the NAs

Before starting the analysis, we need to make some preprocessing on our data. Let us start by loading it into memory and listing the names of the columns.

```
data <- read.csv("diabetes.csv")
colnames(data)
```

```
## [1] "Pregnancies"          "Glucose"
## [3] "BloodPressure"        "SkinThickness"
## [5] "Insulin"              "BMI"
## [7] "DiabetesPedigreeFunction" "Age"
## [9] "Outcome"
```

Of these, our target variable is `Outcome`, which has two levels. For convenience, we transform it into a factor variable which R can treat accordingly.

```
data$Outcome <- factor(data$Outcome, c(0, 1), c("Negative", "Positive"))
```

Now we need to address a particularity of the chosen data: not-a-number (NaN) instances are encoded as zeros in variables where that value is impossible¹. These are:

- Glucose
- BloodPressure
- SkinThickness
- Insulin
- BMI

In order for us to later treat them correctly, we need to manually change them to the existing NA type. As we do so, we record the number of NaNs instances in each of those variables. For convenience, we define a function.

```
set_nas <- function(data, fields) {
  percentage <- list()
  for (field in fields) {
    data[[field]][data[[field]] == 0] <- NA
    percentage[[field]] <- 100 * sum(is.na(data[[field]])) / nrow(data)
  }
}
```

¹Remember we are dealing with medical data, not with artificial one. Hence, there are constraints on the values a variable may take.


```

    }
    return(list(data = data, percentage = percentage))
}

# Correctly label NaNs
na_fields <- c("Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI")
data_na <- set_nas(data, na_fields)
data <- data_na$data
percentages <- data_na$percentage

# Visualize them
vis_miss(data)

```



Now the next logical step is to impute those NaN values. We do have some concern about the imputation of the “Insulin” variable which is almost half-filled of NaNs. Nevertheless, we decide to impute it. On the followed strategy, we use the “Predictive Mean Matching Imputation” (PMM in short) as it behaves much more robustly than naive mean or median imputations.

```
data_im <- mice(data, m = 1, method = "pmm")
```

```
##
## iter imp variable
## 1 1 Glucose BloodPressure SkinThickness Insulin BMI
## 2 1 Glucose BloodPressure SkinThickness Insulin BMI
## 3 1 Glucose BloodPressure SkinThickness Insulin BMI
## 4 1 Glucose BloodPressure SkinThickness Insulin BMI
## 5 1 Glucose BloodPressure SkinThickness Insulin BMI

data <- complete(data_im)
```

Visualization of the data

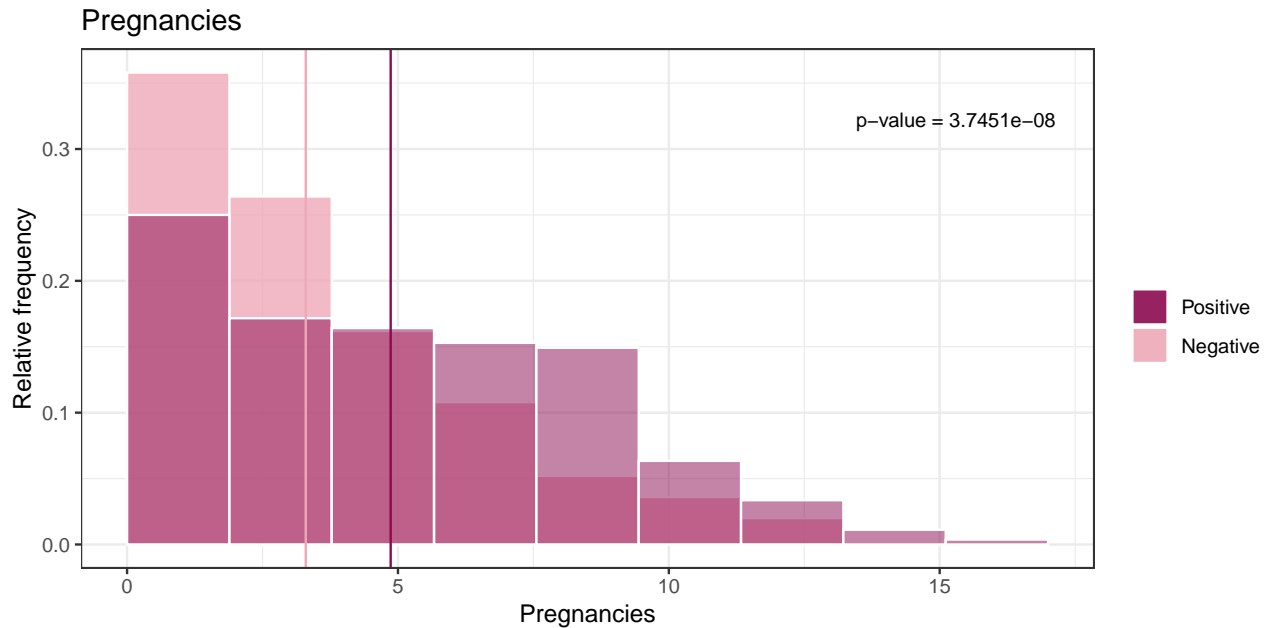
Before we proceed any further, we are going to describe our data. First, we look individually to each of the attributes, we see both visually and with a two-sample Wilcoxon-test whether there is significant difference between the two groups (having or not diabetes), and if some transformation may be desirable to ensure normality compatibility. For that purpose, we define the following function.

```
histogram_by_groups <- function(data, var, label = NULL) {
  stat_t <- wilcox.test(as.formula(paste(var, "~ Outcome")), data)
  data0 <- data[data$Outcome == "Negative", ]
  data1 <- data[data$Outcome == "Positive", ]
  if (is.null(label)) {
    label <- var
  }
  p <- ggplot(data0, aes(x = eval(parse(text = var)))) +
    geom_histogram(
      aes(y = after_stat(count / sum(count)), fill = "Negative"),
      bins = 10, colour = "white", alpha = 0.8, boundary = 0
    ) +
    geom_histogram(data = data1,
      aes(
        x = eval(parse(text = var)),
        y = after_stat(count / sum(count)), fill = "Positive"
      ),
      bins = 10, colour = "white",
      alpha = 0.5, boundary = 0, inherit.aes = FALSE) +
    theme_bw() +
    scale_fill_manual(
      name = "",
      breaks = c("Positive", "Negative"),
      values = c("Positive" = "deeppink4", "Negative" = "pink2")
    ) +
    xlab(label) + ylab("Relative frequency") + ggtitle(label) +
    geom_vline(xintercept = mean(data1[[var]]), colour = "deeppink4") +
    geom_vline(xintercept = mean(data0[[var]]), colour = "pink2")
  p + annotate(
    "text",
    x = 0.9 * max(data1[[var]]),
    y = 0.9 * max(ggplot_build(p)$data[[1]]["y"]),
    label = sprintf("p-value = %.4e", stat_t$p.value),
    size = 3
  )
}
```

Let us start with the number of pregnancies. We can see that people who have diabetes have had more pregnancies than those who do not have diabetes. We see that it seems to be somewhat based on the p-value alone. We also note it exhibits a heavily right-skewed behaviour. As such, a logarithmic transformation would make sense to get a distribution more compatible with the normal one. Nonetheless, a problem here arises in dealing with the null values ². For that purpose we shift the variable by one unit. As a consistency measure, we will apply this shift to all variables we log-transform.


²Remember the domain of the logarithmic function is $(0, \infty)$.

```
histogram_by_groups(data, "Pregnancies")
```



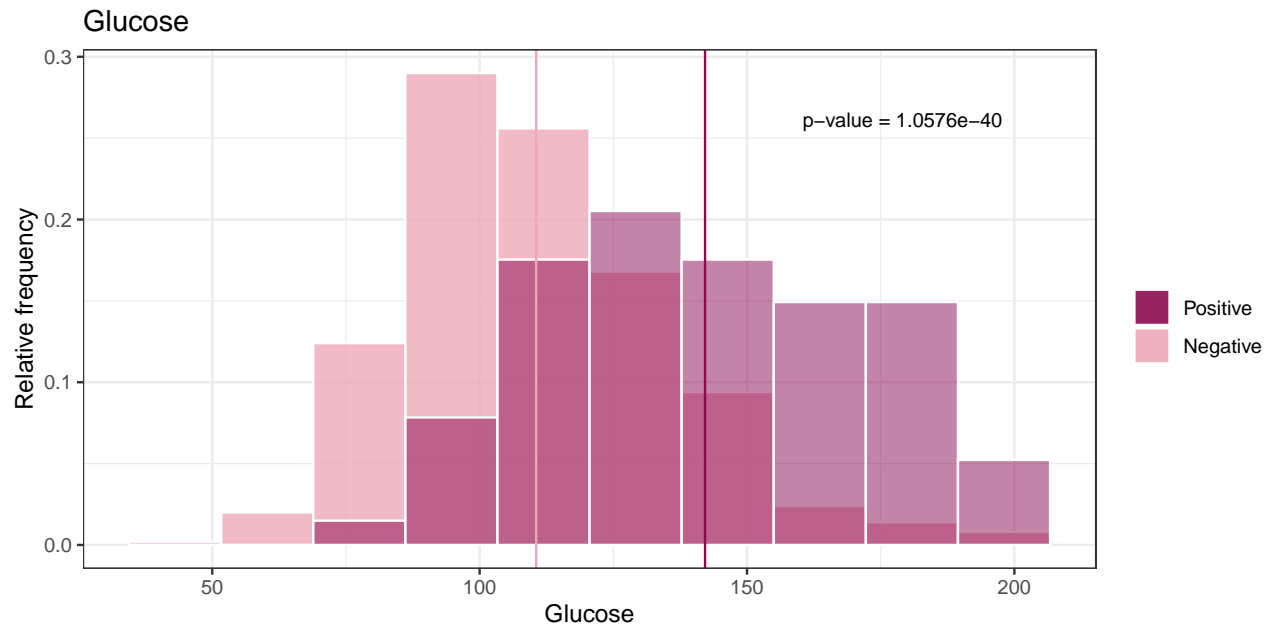
The change as we see does help in obtaining a more centered distribution with which we may better apply the posterior analysis.

```
data$LogPregnancies <- log(data$Pregnancies + 1)
histogram_by_groups(data, "LogPregnancies")
```




The next variable to visualize is the glucose. People with diabetes exhibit higher glucose values. The p-value is very small which indicates a highly significant difference between the two groups. We also observe that the distribution is already well-centered and resembles a normal distribution. Hence, we decide not to transform the data.

```
histogram_by_groups(data, "Glucose")
```



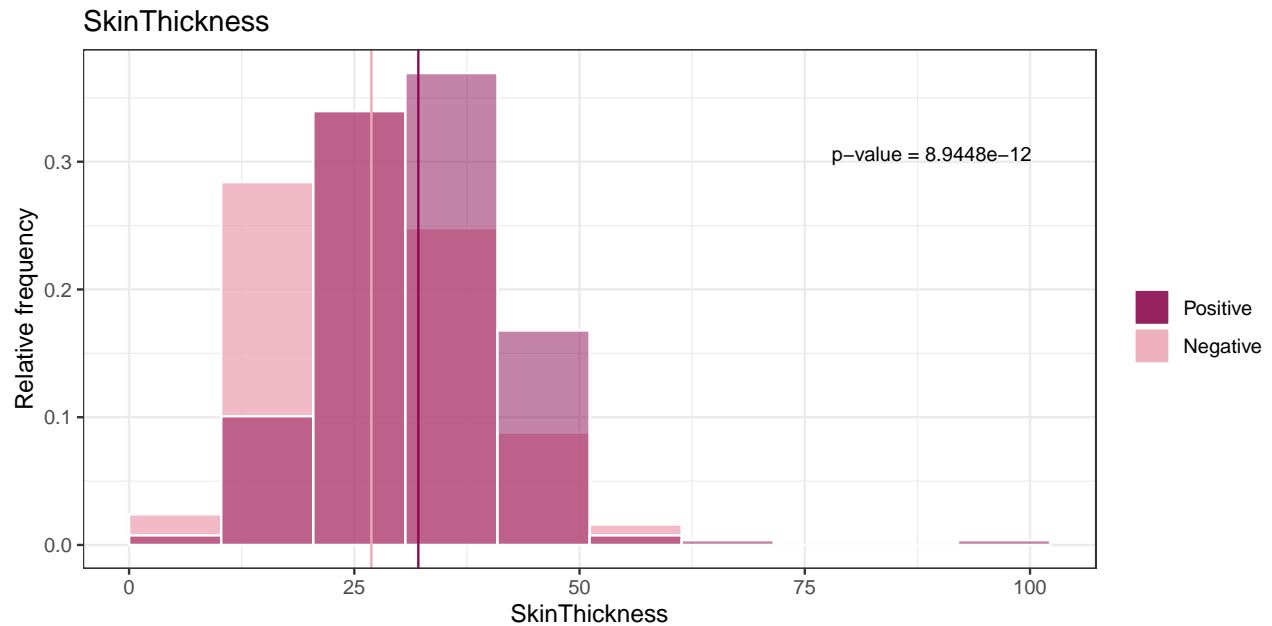
We move onwards to blood pressure. In this case, although the distributions appear as normal, there does not seem to exist a highly significant difference between the groups in contrast to what the p-value states, more so compared with the glucose variable. Nonetheless, there seem to be an slight indication of higher blood pressure for people with diabetes.

```
histogram_by_groups(data, "BloodPressure")
```




Now skin thickness is interesting, because the distributions are visually to those of the blood pressure but the presence of outliers is appreciable. We will later deal with those but for now let us maintain this variable as it is. Note that the population with diabetes appear to exhibit a thicker skin.

```
histogram_by_groups(data, "SkinThickness")
```




Insulin, as we may have expected from the name of the property itself, appears to be a relevant. The median of the distributions does indeed seem to differ, with the one for the diabetes population being slightly higher. It is also right-skewed, so we decide to log-transform it.

```
histogram_by_groups(data, "Insulin")
```




It does appear that the transformation improves the the symmetrization of the data, although some left-skewness appears. We will later see if outlier detection get to target those values or not.

```
data$LogInsulin <- log(data$Insulin + 1)
histogram_by_groups(data, "LogInsulin")
```




Body Mass Index (BMI) again exhibits this tendency of leaning towards a more right-skewed distribution. It does also follow the tendency of being slightly higher for people with diabetes. As such we log-transform to try and get a more normalized variable.

```
histogram_by_groups(data, "BMI")
```




As we may visually judge, it is the case that the log transformation centers the data and provides a more normal-like distribution. There is some presence of seemingly outlying points to the right.

```
data$LogBMI <- log(data$BMI + 1)
histogram_by_groups(data, "LogBMI")
```




The Diabetes Pedigree Function (DPF) is the again a flagrant right-skewed. It again has higher values for the positive set. This is to be expected from the definition of this very function as a risk indication for diabetes. Let us try to log-transform it.

```
histogram_by_groups(data, "DiabetesPedigreeFunction", "DPF")
```




The improvement is highly noticeable. We retain thus this transformed variable.

```
data$LogDPF <- log(data$DiabetesPedigreeFunction + 1)
histogram_by_groups(data, "LogDPF")
```

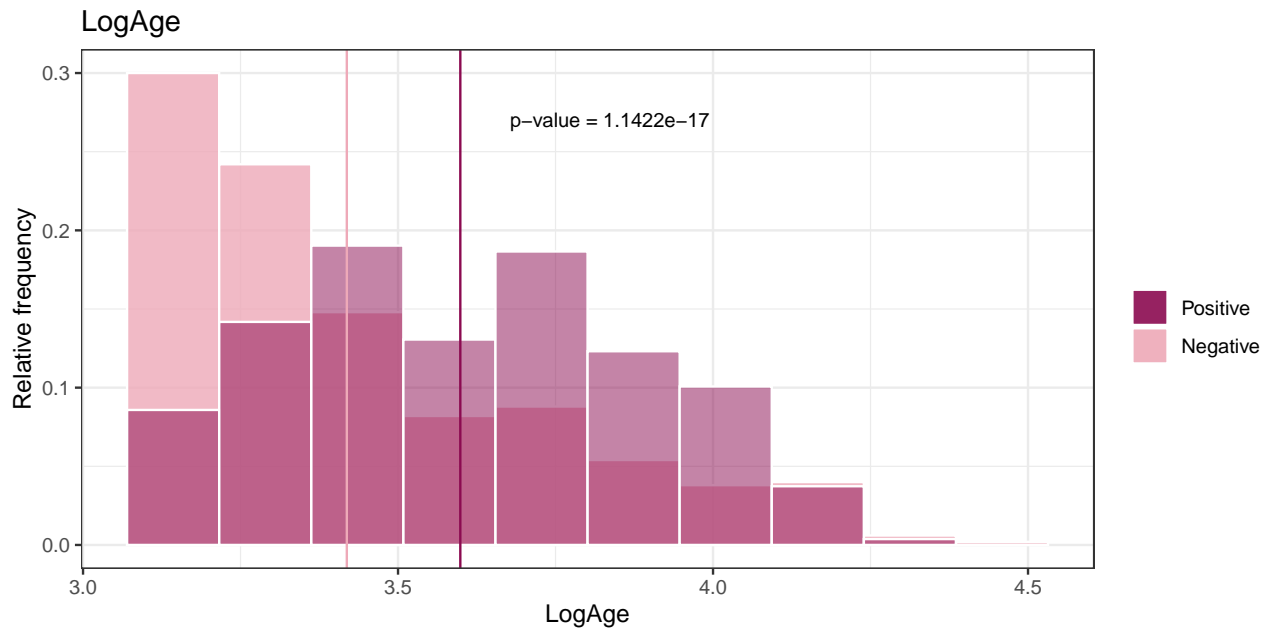
We may note that young people, as with other illnesses have less tendency to suffer diabetes than the elders. The age is expected to exhibit a right-skewed distribution, as is indeed the case. In an attempt to improve the symmetry, we once again use logarithms to transform the variable.

```
histogram_by_groups(data, "Age")
```



The transformation does help although not by much. This is a common problem of the age variable. We keep the transformation anyway as it does seem to help with the symmetry for the positive group.


```
data$LogAge <- log(data$Age + 1)
histogram_by_groups(data, "LogAge")
```



```
# Some convenience
df <- subset(data, select = -c(Pregnancies, Insulin, BMI, DiabetesPedigreeFunction, Age))
df0 <- df[df$Outcome == "Negative", ]
df1 <- df[df$Outcome == "Positive", ]
xnames <- names(df)[! names(df) %in% c("Outcome")]
```

Now, let's take a look at some multivariate plots. We'll begin by inspecting the Parallel Coordinate Plot:

```
par(las = 2)
parcoord(df[xnames], col = c("pink2", "deeppink4"))
legend("topright", legend = c("Negative", "Positive"),
      col = c("pink2", "deeppink4"), lty = 1, lwd = 2)
```




It seems that, overall, the positive lines are over negative ones. This is most notable on the Glucose and log transformed BMI. They are two of the most significant features according to the p-value.

The Andrew's plot is the following:

```
andrewsplot(as.matrix(df[xnames]), df$Outcome, style = "cart")
legend("topright", legend = c("Negative", "Positive"),
       col = c("black", "red"), lty = 1, lwd = 2)
```

Andrews' Curves



Again, we see that the two groups are different. The group of people who have diabetes tend to have more volatile curves, reaching higher and lower values along the curve.

Multivariate characteristics and outlier identification

We are going to use a multivariate approach to identifying the outliers in our data. In the process, we will need to compute the mean and covariance.

We start then by finding the mean vector and the covariance matrix. In order to reduce the sensitivity to outliers in this computation, we use a robust estimation using the “Fast MCD” (Minimum Covariance Determinant) estimator. We set main parameter, α , which determines the percentage of the data to use, at 0.85.

```
our_corrplot <- function(cov_mat) {
  colnames(cov_mat) <- c("Log\nPregnancies",
                        "Glucose",
                        "Blood\nPressure",
                        "Skin\nThickness",
                        "LogInsulin",
                        "LogBMI",
                        "LogDPF",
                        "LogAge")
  corrplot.mixed(cov2cor(cov_mat), lower = "number", upper = "color",
                 diag = "n", tl.col = "black", tl.cex = 0.65,
                 lower.col = "black")
}

mcd_est <- CovMcd(df[xnames], alpha = 0.85, nsamp = "deterministic")
```

```
## [1] "The mean vector is:"
```

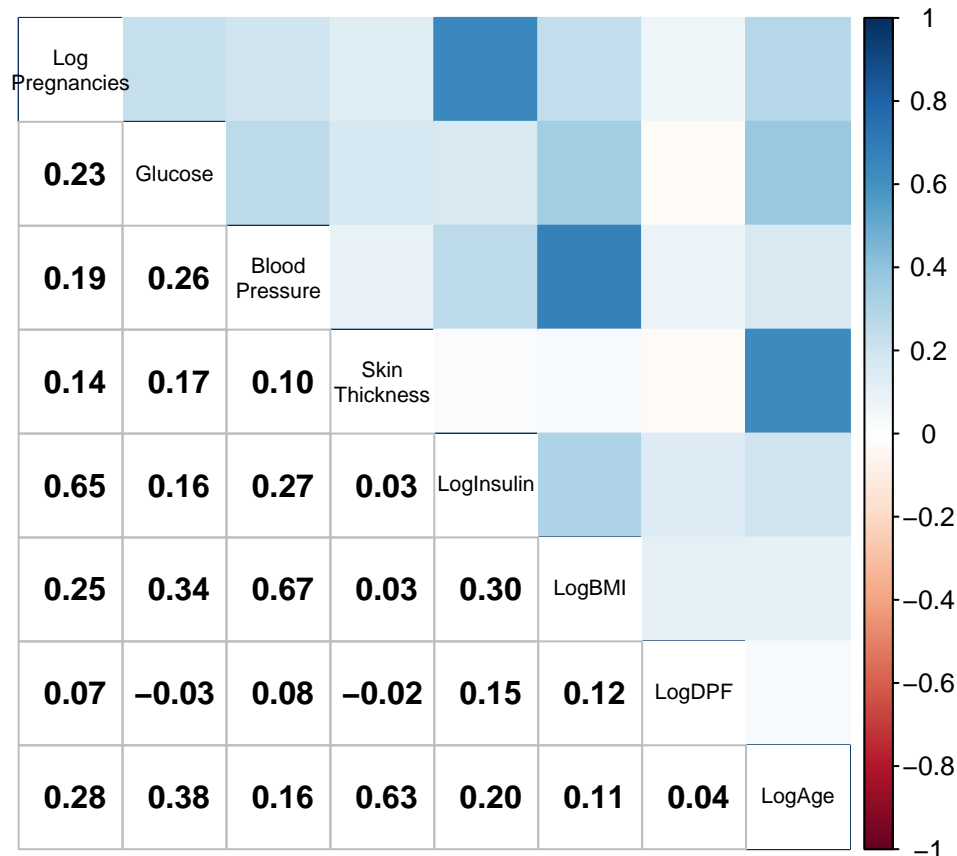
```
##      Glucose  BloodPressure  SkinThickness  LogPregnancies  LogInsulin
##  120.4547983    71.9860918    28.3810848     1.3402737     4.7230964
```

```
##          LogBMI          LogDPF          LogAge
##      3.4853699      0.3514095      3.4677968

## [1] "The covariance matrix is:"

##          Glucose BloodPressure SkinThickness LogPregnancies LogInsulin
## Glucose      949.5595      85.7752      60.9537      3.2722      14.2651
## BloodPressure 85.7752     142.7851      32.1056      1.6069      1.3763
## SkinThickness 60.9537      32.1056     105.0266      0.7861      1.9485
## LogPregnancies 3.2722      1.6069      0.7861      0.6162      0.0143
## LogInsulin    14.2651      1.3763      1.9485      0.0143      0.5123
## LogBMI        1.5692      0.8434      1.4231      0.0055      0.0444
## LogDPF        0.4139     -0.0651      0.1557     -0.0035      0.0200
## LogAge        2.7784      1.4462      0.5345      0.1587      0.0461
##          LogBMI LogDPF LogAge
## Glucose      1.5692 0.4139 2.7784
## BloodPressure 0.8434 -0.0651 1.4462
## SkinThickness 1.4231 0.1557 0.5345
## LogPregnancies 0.0055 -0.0035 0.1587
## LogInsulin    0.0444 0.0200 0.0461
## LogBMI        0.0427 0.0046 0.0072
## LogDPF        0.0046 0.0353 0.0022
## LogAge        0.0072 0.0022 0.1019
```

```
our_corrplot(mcd_est$cov)
```



It is interesting to segregate by the class of the Outcome variable. This way, we can both, get the mean vector and covariance matrix for each of the classes, and the outliers for both sets.

```

mcd_neg <- CovMcd(df0[xnames], alpha = 0.85, nsamp = "deterministic")
mcd_pos <- CovMcd(df1[xnames], alpha = 0.85, nsamp = "deterministic")

## [1] "### Negative class ###"
## [1] "The mean vector is:"

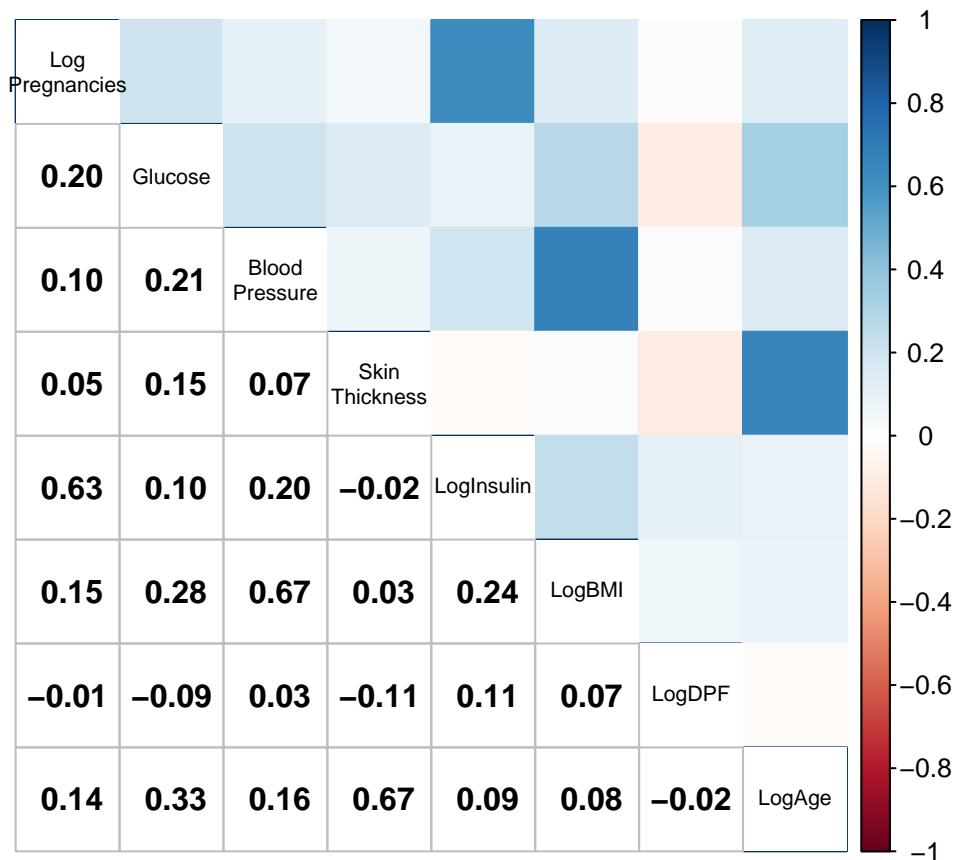
##      Glucose  BloodPressure  SkinThickness  LogPregnancies  LogInsulin
##  108.1916300    70.0242291    26.6365639      1.2186616      4.5363303
##      LogBMI      LogDPF      LogAge
##    3.4390348    0.3232806    3.3844990

## [1] "The covariance matrix is:"

##      Glucose  BloodPressure  SkinThickness  LogPregnancies  LogInsulin
## Glucose      535.7768      53.3925      24.9768      0.9621      10.0637
## BloodPressure  53.3925     132.1672      25.3349      1.2742      0.7898
## SkinThickness  24.9768     25.3349     110.6200      0.5693      1.4467
## LogPregnancies  0.9621      1.2742      0.5693      0.5732     -0.0112
## LogInsulin     10.0637      0.7898      1.4467     -0.0112      0.4821
## LogBMI          0.7120      0.6712      1.4900      0.0041      0.0355
## LogDPF          -0.0496     -0.1801      0.0490     -0.0139      0.0135
## LogAge          0.9492      1.1183      0.4847      0.1485      0.0188
##      LogBMI  LogDPF  LogAge
## Glucose      0.7120 -0.0496 0.9492
## BloodPressure 0.6712 -0.1801 1.1183
## SkinThickness 1.4900 0.0490 0.4847
## LogPregnancies 0.0041 -0.0139 0.1485
## LogInsulin     0.0355 0.0135 0.0188
## LogBMI          0.0447 0.0025 0.0053
## LogDPF          0.0025 0.0294 -0.0013
## LogAge          0.0053 -0.0013 0.0866

our_corrplot(mcd_neg$cov)

```



```
## [1] "### Positive class ###"
```

```
## [1] "The mean vector is:"
```

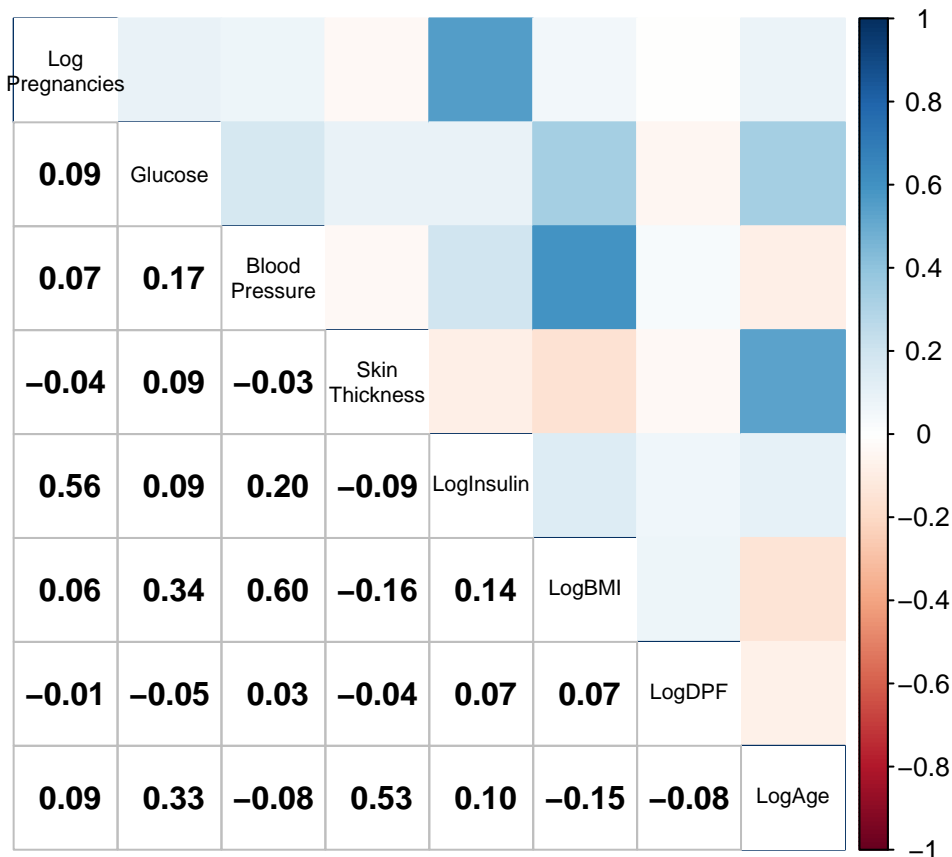
```
##      Glucose  BloodPressure  SkinThickness  LogPregnancies  LogInsulin
##  142.2148438    75.5351562    31.7304688    1.5230344    5.0556806
##      LogBMI      LogDPF      LogAge
##    3.5763371    0.4027727    3.5985651
```

```
## [1] "The covariance matrix is:"
```

```
##      Glucose  BloodPressure  SkinThickness  LogPregnancies  LogInsulin
## Glucose      931.9767      34.5571      21.1150      -1.0132      11.1108
## BloodPressure 34.5571      152.7967      19.9239      0.9369      0.7668
## SkinThickness 21.1150      19.9239      86.7544      -0.2417      1.2014
## LogPregnancies -1.0132      0.9369      -0.2417      0.6889      -0.0464
## LogInsulin    11.1108      0.7668      1.2014      -0.0464      0.4275
## LogBMI        0.2897      0.7166      0.9532      -0.0222      0.0159
## LogDPF        -0.0499      -0.1220      0.0606      -0.0064      0.0091
## LogAge        0.7751      1.1984      -0.2209      0.1289      0.0199
##      LogBMI  LogDPF  LogAge
## Glucose    0.2897 -0.0499 0.7751
## BloodPressure 0.7166 -0.1220 1.1984
## SkinThickness 0.9532 0.0606 -0.2209
## LogPregnancies -0.0222 -0.0064 0.1289
## LogInsulin    0.0159 0.0091 0.0199
## LogBMI        0.0294 0.0026 -0.0075
## LogDPF        0.0026 0.0430 -0.0047
```

```
## LogAge          -0.0075 -0.0047  0.0854
```

```
our_corrplot(mcd_pos$cov)
```



Let us focus first on the mean vectors. The here observed particularities are nothing new, as they were already present on above histograms:

- Higher number of pregnancies seems to increase the chances on developing diabetes.
- The glucose and insulin levels of the positive population is higher in contrast to the negative one.

Looking now at the representations for the covariance matrices, the major changes are that the correlation between skin thickness and diabetes pedigree function is lower in the group who do not have diabetes than in the one having the disease. Moreover, the correlation between BMI³ and age is positive for the sane group while negative on the positive one.

We now search for outliers. The idea is to use Mahalanobis distance. As we suppose that our data $X \sim \mathcal{N}(\mu, \Sigma)$, we have $D_M(x, \mu)^2 \sim \chi_p^2$, with D_M the Mahalanobis distance. Hence we may set our outlier criteria as

$$D_M(x, \mu)^2 > \chi_{p, 0.95^{1/n}}^2, \quad (1)$$

the 0.95th quantile of the χ_p^2 distribution. We then drop these outliers.

```
p <- length(xnames)
n0 <- nrow(df0)
n1 <- nrow(df1)
df0_clean <- df0[mcd_neg$mah < qchisq(0.95^(1 / n0), p), ]
df1_clean <- df1[mcd_pos$mah < qchisq(0.95^(1 / n1), p), ]
df_clean <- rbind(df0_clean, df1_clean)
```

³Remember this is the logarithm.

```
## [1] "The negative set contained 6 outliers."
```

```
## [1] "The positive set contained 4 outliers."
```

As a final summary of this section, we perform a plot in which the histograms of different populations are observed, as well as scatterplots of pairs of variables and correlations.

```
ggpairs(df_clean, aes(color = Outcome), legend = 1,
        columns = xnames,
        diag = list(continuous = "barDiag")
    ) +
    theme(legend.position = "bottom") +
    scale_fill_manual(values = c("pink", "deeppink4")) +
    scale_color_manual(values = c("pink", "deeppink4")) + labs(fill = "Outcome")
```

